

# Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

# Importing Datasets

In [2]:

```
df=pd.read_csv("2016.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	s
0	2016-11-01 01:00:00	NaN	0.7	NaN	NaN	153.0	77.0	NaN	NaN	NaN	7.0	NaN	NaN	280
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.4	280
2	2016-11-01 01:00:00	5.9	NaN	7.5	NaN	297.0	139.0	NaN	NaN	NaN	NaN	NaN	26.0	280
3	2016-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	113.0	2.0	NaN	NaN	NaN	NaN	NaN	280
4	2016-11-01 01:00:00	NaN	NaN	NaN	NaN	275.0	127.0	2.0	NaN	NaN	18.0	NaN	NaN	280
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
209491	2016-07-01 00:00:00	NaN	0.2	NaN	NaN	2.0	29.0	73.0	NaN	NaN	NaN	NaN	NaN	280
209492	2016-07-01 00:00:00	NaN	0.3	NaN	NaN	1.0	29.0	NaN	36.0	NaN	5.0	NaN	NaN	280
209493	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	1.0	19.0	71.0	NaN	NaN	NaN	NaN	NaN	280
209494	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	6.0	17.0	85.0	NaN	NaN	NaN	NaN	NaN	280
209495	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	2.0	46.0	61.0	34.0	NaN	NaN	NaN	NaN	280

209496 rows × 14 columns

# Data Cleaning and Data Preprocessing

In [3]: `df=df.dropna()`

In [4]: `df.columns`

Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO\_2', 'O\_3', 'PM10', 'PM25', 'SO\_2', 'TCH', 'TOL', 'station'], dtype='object')

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 16932 entries, 1 to 209478
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        16932 non-null   object
1   BEN         16932 non-null   float64
2   CO          16932 non-null   float64
3   EBE         16932 non-null   float64
4   NMHC        16932 non-null   float64
5   NO          16932 non-null   float64
6   NO_2        16932 non-null   float64
7   O_3         16932 non-null   float64
8   PM10        16932 non-null   float64
9   PM25        16932 non-null   float64
10  SO_2        16932 non-null   float64
11  TCH         16932 non-null   float64
12  TOL         16932 non-null   float64
13  station     16932 non-null   int64
dtypes: float64(12), int64(1), object(1)
memory usage: 1.9+ MB
```

In [6]: `data=df[['CO', 'station']]`  
`data`

Out[6]:


	CO	station
1	1.1	28079008
6	0.8	28079024
25	1.0	28079008
30	0.7	28079024
49	0.8	28079008
...	...	...
209430	0.2	28079024
209449	0.4	28079008
209454	0.2	28079024
209473	0.4	28079008
209478	0.2	28079024

16932 rows × 2 columns

## Line chart

```
In [7]: data.plot.line(subplots=True)
```


```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



## Line chart

```
In [8]: data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```




## Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```


```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: data.plot.hist()
```


```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

```
In [12]: data.plot.area()
```


```
Out[12]: <AxesSubplot:>
```



## Box chart

```
In [13]: data.plot.box()
```


```
Out[13]: <AxesSubplot:>
```



## Pie chart

```
In [14]: b.plot.pie(y='station' )
```


```
Out[14]: <AxesSubplot:ylabel='station'>
```



# Scatter chart

```
In [15]: data.plot.scatter(x='CO' ,y='station')
```

Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>



In [16]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 16932 entries, 1 to 209478
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        16932 non-null  object
1   BEN         16932 non-null  float64
2   CO          16932 non-null  float64
3   EBE         16932 non-null  float64
4   NMHC        16932 non-null  float64
5   NO          16932 non-null  float64
6   NO_2        16932 non-null  float64
7   O_3         16932 non-null  float64
8   PM10        16932 non-null  float64
9   PM25        16932 non-null  float64
10  SO_2        16932 non-null  float64
11  TCH         16932 non-null  float64
12  TOL         16932 non-null  float64
13  station     16932 non-null  int64
dtypes: float64(12), int64(1), object(1)
memory usage: 1.9+ MB
```

In [17]:

df.columns

Out[17]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO\_2', 'O\_3', 'PM10', 'PM25', 'SO\_2', 'TCH', 'TOL', 'station'], dtype='object')

In [18]:

df.describe()

Out[18]:

	BEN	CO	EBE	NMHC	NO	NO_2	
count	16932.000000	16932.000000	16932.000000	16932.000000	16932.000000	16932.000000	16932.00
mean	0.537970	0.349941	0.298955	0.099913	20.815734	39.373376	48.11
std	0.599479	0.203807	0.450204	0.079850	40.986063	31.170307	32.56
min	0.100000	0.100000	0.100000	0.000000	1.000000	1.000000	1.00
25%	0.200000	0.200000	0.100000	0.050000	1.000000	14.000000	21.00
50%	0.400000	0.300000	0.200000	0.090000	7.000000	34.000000	46.00
75%	0.700000	0.400000	0.300000	0.120000	23.000000	58.000000	69.00


	BEN	CO	EBE	NMHC	NO	NO_2	NO_2
max	12.300000	4.500000	13.500000	2.210000	829.000000	319.000000	181.00

```
In [19]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2','O_3',
               'PM10','SO_2', 'TCH', 'TOL', 'station']]
```

## EDA AND VISUALIZATION

```
In [20]: sns.pairplot(df1[0:50])
```

```
Out[20]: <seaborn.axisgrid.PairGrid at 0x217d78acc70>
```




```
In [21]: sns.distplot(df1['station'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar f




lexibility) or `histplot` (an axes-level function for histograms).  
 warnings.warn(msg, FutureWarning)

Out[21]: <AxesSubplot:xlabel='station', ylabel='Density'>



In [22]: `sns.heatmap(df1.corr())`

Out[22]: <AxesSubplot:>



## TO TRAIN THE MODEL AND MODEL BUILDING

In [23]: `x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',  
           'PM10', 'SO_2', 'TCH', 'TOL']]  
 y=df['station']`

In [24]: `from sklearn.model_selection import train_test_split  
 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)`

## Linear Regression

```
In [25]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[25]: LinearRegression()

```
In [26]: lr.intercept_
```

Out[26]: 28079040.05833833

```
In [27]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```


Out[27]:

	Co-efficient
<b>BEN</b>	1.799305
<b>CO</b>	5.866820
<b>EBE</b>	0.210592
<b>NMHC</b>	3.652970
<b>NO_2</b>	-0.067171
<b>O_3</b>	-0.028079
<b>PM10</b>	0.024436
<b>SO_2</b>	-0.844531
<b>TCH</b>	-12.959269
<b>TOL</b>	0.258683

	Co-efficient
<b>BEN</b>	1.799305
<b>CO</b>	5.866820
<b>EBE</b>	0.210592
<b>NMHC</b>	3.652970
<b>NO_2</b>	-0.067171
<b>O_3</b>	-0.028079
<b>PM10</b>	0.024436
<b>SO_2</b>	-0.844531
<b>TCH</b>	-12.959269
<b>TOL</b>	0.258683

```
In [28]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[28]: <matplotlib.collections.PathCollection at 0x217e0d9bee0>



## ACCURACY

```
In [29]: lr.score(x_test,y_test)
```

```
Out[29]: 0.8011739192797894
```

```
In [30]: lr.score(x_train,y_train)
```

```
Out[30]: 0.7916349423598931
```

## Ridge and Lasso

```
In [31]: from sklearn.linear_model import Ridge,Lasso
```

```
In [32]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
Out[32]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [33]: rr.score(x_test,y_test)
```

```
Out[33]: 0.8004456471536664
```

```
In [34]: rr.score(x_train,y_train)
```

```
Out[34]: 0.7915061315853162
```

```
In [35]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[35]: Lasso(alpha=10)
```

```
In [36]: la.score(x_train,y_train)
```

```
Out[36]: 0.6160284065614383
```

## Accuracy(Lasso)

```
In [37]: la.score(x_test,y_test)
```

```
Out[37]: 0.6262405487220697
```

## Elastic Net

```
In [38]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(x_train,y_train)
```

Out[38]: ElasticNet()

```
In [39]: en.coef_
```

Out[39]: array([ 0. , 0. , 0. , -0. , -0.07934312,  
 -0.02660259, 0.02377698, -0.85401015, -0. , 0.2541252 ])

```
In [40]: en.intercept_
```

Out[40]: 28079025.817570496

```
In [41]: prediction=en.predict(x_test)
```

```
In [42]: en.score(x_test,y_test)
```

Out[42]: 0.686409247096778

## Evaluation Metrics

```
In [43]: from sklearn import metrics
         print(metrics.mean_absolute_error(y_test,prediction))
         print(metrics.mean_squared_error(y_test,prediction))
         print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

3.5275189590574367  
20.069655755308172  
4.4799169362063145

## Logistic Regression

```
In [44]: from sklearn.linear_model import LogisticRegression
```

```
In [45]: feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2','O_3',  
                           'PM10','SO_2', 'TCH', 'TOL']]
         target_vector=df['station']
```

```
In [46]: feature_matrix.shape
```

Out[46]: (16932, 10)

```
In [47]: target_vector.shape
```

Out[47]: (16932,)

```
In [48]: from sklearn.preprocessing import StandardScaler
```

```
In [49]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [50]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[50]: LogisticRegression(max_iter=10000)
```

```
In [51]: observation=[[1,2,3,4,5,6,7,8,9,10]]
```

```
In [52]: prediction=logr.predict(observation)
print(prediction)
```

```
[28079008]
```

```
In [53]: logr.classes_
```

```
Out[53]: array([28079008, 28079024], dtype=int64)
```

```
In [54]: logr.score(fs,target_vector)
```

```
Out[54]: 0.9923812898653437
```

```
In [55]: logr.predict_proba(observation)[0][0]
```

```
Out[55]: 1.0
```

```
In [56]: logr.predict_proba(observation)
```

```
Out[56]: array([[1.0000000e+00, 1.6336121e-46]])
```

## Random Forest

```
In [57]: from sklearn.ensemble import RandomForestClassifier
```

```
In [58]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[58]: RandomForestClassifier()
```

```
In [59]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]
}
```

```
In [60]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc, param_grid=parameters, cv=2, scoring="accuracy")
grid_search.fit(x_train, y_train)
```

```
Out[60]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

```
In [61]: grid_search.best_score_
```

```
Out[61]: 0.9946000674991562
```

```
In [62]: rfc_best = grid_search.best_estimator_
```

```
In [63]: from sklearn.tree import plot_tree

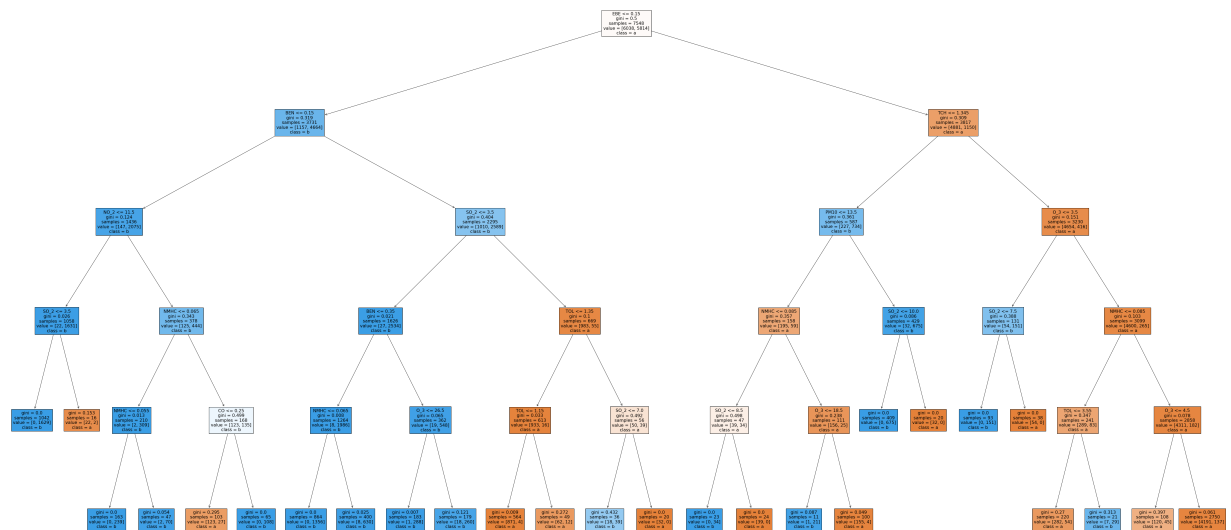
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b', 'c'],
```

```
Out[63]: [Text(2266.1632653061224, 1993.2, 'EBE <= 0.15\nngini = 0.5\nnsamples = 7548\nvalue =
[6038, 5814]\nnclass = a'),
Text(1070.4489795918369, 1630.8000000000002, 'BEN <= 0.15\nngini = 0.319\nnsamples =
3731\nvalue = [1157, 4664]\nnclass = b'),
Text(409.9591836734694, 1268.4, 'NO_2 <= 11.5\nngini = 0.124\nnsamples = 1436\nvalue
= [147, 2075]\nnclass = b'),
Text(182.20408163265307, 906.0, 'SO_2 <= 3.5\nngini = 0.026\nnsamples = 1058\nvalue =
[22, 1631]\nnclass = b'),
Text(91.10204081632654, 543.5999999999999, 'gini = 0.0\nnsamples = 1042\nvalue = [0,
1629]\nnclass = b'),
Text(273.30612244897964, 543.5999999999999, 'gini = 0.153\nnsamples = 16\nvalue = [2
2, 2]\nnclass = a'),
Text(637.7142857142858, 906.0, 'NMHC <= 0.065\nngini = 0.343\nnsamples = 378\nvalue =
[125, 444]\nnclass = b'),
Text(455.51020408163265, 543.5999999999999, 'NMHC <= 0.055\nngini = 0.013\nnsamples =
210\nvalue = [2, 309]\nnclass = b'),
Text(364.40816326530614, 181.19999999999982, 'gini = 0.0\nnsamples = 163\nvalue =
[0, 239]\nnclass = b'),
Text(546.6122448979593, 181.19999999999982, 'gini = 0.054\nnsamples = 47\nvalue =
[2, 70]\nnclass = b'),
Text(819.9183673469388, 543.5999999999999, 'CO <= 0.25\nngini = 0.499\nnsamples = 168
\nvalue = [123, 135]\nnclass = b'),
Text(728.8163265306123, 181.19999999999982, 'gini = 0.295\nnsamples = 103\nvalue =
[123, 27]\nnclass = a'),
Text(911.0204081632653, 181.19999999999982, 'gini = 0.0\nnsamples = 65\nvalue = [0,
108]\nnclass = b'),
Text(1730.938775510204, 1268.4, 'SO_2 <= 3.5\nngini = 0.404\nnsamples = 2295\nvalue =
[1010, 2589]\nnclass = b'),
Text(1366.530612244898, 906.0, 'BEN <= 0.35\nngini = 0.021\nnsamples = 1626\nvalue =
[27, 2534]\nnclass = b'),
Text(1184.326530612245, 543.5999999999999, 'NMHC <= 0.065\nngini = 0.008\nnsamples =
1264\nvalue = [8, 1986]\nnclass = b'),
Text(1093.2244897959185, 181.19999999999982, 'gini = 0.0\nnsamples = 864\nvalue =
[0, 1356]\nnclass = b'),
Text(1275.4285714285716, 181.19999999999982, 'gini = 0.025\nnsamples = 400\nvalue =
[8, 630]\nnclass = b'),
Text(1548.734693877551, 543.5999999999999, 'O_3 <= 26.5\nngini = 0.065\nnsamples = 36
2\nvalue = [19, 548]\nnclass = b'),
Text(1457.6326530612246, 181.19999999999982, 'gini = 0.007\nnsamples = 183\nvalue =
[1, 288]\nnclass = b'),
Text(1639.8367346938776, 181.19999999999982, 'gini = 0.121\nnsamples = 179\nvalue =
```

```

[18, 260]\n\nclass = b'),
Text(2095.3469387755104, 906.0, 'TOL <= 1.35\n\ngini = 0.1\n\nsamples = 669\n\nvalue = [9
83, 55]\n\nclass = a'),
Text(1913.1428571428573, 543.5999999999999, 'TOL <= 1.15\n\ngini = 0.033\n\nsamples = 6
13\n\nvalue = [933, 16]\n\nclass = a'),
Text(1822.0408163265306, 181.19999999999982, 'gini = 0.009\n\nsamples = 564\n\nvalue =
[871, 4]\n\nclass = a'),
Text(2004.2448979591838, 181.19999999999982, 'gini = 0.272\n\nsamples = 49\n\nvalue =
[62, 12]\n\nclass = a'),
Text(2277.5510204081634, 543.5999999999999, 'SO_2 <= 7.0\n\ngini = 0.492\n\nsamples = 5
6\n\nvalue = [50, 39]\n\nclass = a'),
Text(2186.448979591837, 181.19999999999982, 'gini = 0.432\n\nsamples = 36\n\nvalue = [1
8, 39]\n\nclass = b'),
Text(2368.65306122449, 181.19999999999982, 'gini = 0.0\n\nsamples = 20\n\nvalue = [32,
0]\n\nclass = a'),
Text(3461.877551020408, 1630.8000000000002, 'TCH <= 1.345\n\ngini = 0.309\n\nsamples =
3817\n\nvalue = [4881, 1150]\n\nclass = a'),
Text(3051.918367346939, 1268.4, 'PM10 <= 13.5\n\ngini = 0.361\n\nsamples = 587\n\nvalue =
[227, 734]\n\nclass = b'),
Text(2824.1632653061224, 906.0, 'NMHC <= 0.085\n\ngini = 0.357\n\nsamples = 158\n\nvalue
= [195, 59]\n\nclass = a'),
Text(2641.9591836734694, 543.5999999999999, 'SO_2 <= 8.5\n\ngini = 0.498\n\nsamples = 4
7\n\nvalue = [39, 34]\n\nclass = a'),
Text(2550.857142857143, 181.19999999999982, 'gini = 0.0\n\nsamples = 23\n\nvalue = [0,
34]\n\nclass = b'),
Text(2733.061224489796, 181.19999999999982, 'gini = 0.0\n\nsamples = 24\n\nvalue = [39,
0]\n\nclass = a'),
Text(3006.367346938776, 543.5999999999999, 'O_3 <= 18.5\n\ngini = 0.238\n\nsamples = 11
1\n\nvalue = [156, 25]\n\nclass = a'),
Text(2915.265306122449, 181.19999999999982, 'gini = 0.087\n\nsamples = 11\n\nvalue =
[1, 21]\n\nclass = b'),
Text(3097.469387755102, 181.19999999999982, 'gini = 0.049\n\nsamples = 100\n\nvalue =
[155, 4]\n\nclass = a'),
Text(3279.673469387755, 906.0, 'SO_2 <= 10.0\n\ngini = 0.086\n\nsamples = 429\n\nvalue =
[32, 675]\n\nclass = b'),
Text(3188.571428571429, 543.5999999999999, 'gini = 0.0\n\nsamples = 409\n\nvalue = [0,
675]\n\nclass = b'),
Text(3370.775510204082, 543.5999999999999, 'gini = 0.0\n\nsamples = 20\n\nvalue = [32,
0]\n\nclass = a'),
Text(3871.8367346938776, 1268.4, 'O_3 <= 3.5\n\ngini = 0.151\n\nsamples = 3230\n\nvalue =
[4654, 416]\n\nclass = a'),
Text(3644.081632653061, 906.0, 'SO_2 <= 7.5\n\ngini = 0.388\n\nsamples = 131\n\nvalue =
[54, 151]\n\nclass = b'),
Text(3552.979591836735, 543.5999999999999, 'gini = 0.0\n\nsamples = 93\n\nvalue = [0, 1
51]\n\nclass = b'),
Text(3735.183673469388, 543.5999999999999, 'gini = 0.0\n\nsamples = 38\n\nvalue = [54,
0]\n\nclass = a'),
Text(4099.591836734694, 906.0, 'NMHC <= 0.085\n\ngini = 0.103\n\nsamples = 3099\n\nvalue
= [4600, 265]\n\nclass = a'),
Text(3917.387755102041, 543.5999999999999, 'TOL <= 3.55\n\ngini = 0.347\n\nsamples = 24
1\n\nvalue = [289, 83]\n\nclass = a'),
Text(3826.2857142857147, 181.19999999999982, 'gini = 0.27\n\nsamples = 220\n\nvalue =
[282, 54]\n\nclass = a'),
Text(4008.4897959183677, 181.19999999999982, 'gini = 0.313\n\nsamples = 21\n\nvalue =
[7, 29]\n\nclass = b'),
Text(4281.795918367347, 543.5999999999999, 'O_3 <= 4.5\n\ngini = 0.078\n\nsamples = 285
8\n\nvalue = [4311, 182]\n\nclass = a'),
Text(4190.693877551021, 181.19999999999982, 'gini = 0.397\n\nsamples = 108\n\nvalue =
[120, 45]\n\nclass = a'),
Text(4372.897959183674, 181.19999999999982, 'gini = 0.061\n\nsamples = 2750\n\nvalue =
[4191, 137]\n\nclass = a')]

```



## Conclusion

## Scores

## Linear Regression

In [64]: `lr.score(x_test,y_test)`

Out[64]: 0.8011739192797894

In [65]: `lr.score(x_train,y_train)`

Out[65]: 0.7916349423598931

## Lasso

In [66]: `la.score(x_test,y_test)`

Out[66]: 0.6262405487220697

## Ridge

In [67]: `rr.score(x_test,y_test)`

Out[67]: 0.8004456471536664

In [68]: `rr.score(x_train,y_train)`

Out[68]: 0.7915061315853162



## Elastic Net

```
In [69]: en.score(x_test,y_test)
```

```
Out[69]: 0.686409247096778
```

## Logistic Regression

```
In [70]: logr.score(fs,target_vector)
```

```
Out[70]: 0.9923812898653437
```

## Random Forest

```
In [71]: grid_search.best_score_
```

```
Out[71]: 0.9946000674991562
```

From the above data, we can conclude that random forest regression and logistic regression is preferable to other regression types

```
In [ ]:
```