

Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

In [2]:

```
df=pd.read_csv("2017.csv")
df
```

Out[2]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH
0	2017-06-01 01:00:00	NaN	NaN	0.3	NaN	NaN	4.0	38.0	NaN	NaN	NaN	NaN	5.0	NaN
1	2017-06-01 01:00:00	0.6	NaN	0.3	0.4	0.08	3.0	39.0	NaN	71.0	22.0	9.0	7.0	1.4
2	2017-06-01 01:00:00	0.2	NaN	NaN	0.1	NaN	1.0	14.0	NaN	NaN	NaN	NaN	NaN	NaN
3	2017-06-01 01:00:00	NaN	NaN	0.2	NaN	NaN	1.0	9.0	NaN	91.0	NaN	NaN	NaN	NaN
4	2017-06-01 01:00:00	NaN	NaN	NaN	NaN	NaN	1.0	19.0	NaN	69.0	NaN	NaN	2.0	NaN
...
210115	2017-08-01 00:00:00	NaN	NaN	0.2	NaN	NaN	1.0	27.0	NaN	65.0	NaN	NaN	NaN	NaN
210116	2017-08-01 00:00:00	NaN	NaN	0.2	NaN	NaN	1.0	14.0	NaN	NaN	73.0	NaN	7.0	NaN
210117	2017-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	1.0	4.0	NaN	83.0	NaN	NaN	NaN	NaN
210118	2017-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	1.0	11.0	NaN	78.0	NaN	NaN	NaN	NaN
210119	2017-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	1.0	14.0	NaN	77.0	60.0	NaN	NaN	NaN

210120 rows × 16 columns

Data Cleaning and Data Preprocessing

In [3]: `df=df.dropna()`

In [4]: `df.columns`

Out[4]: Index(['date', 'BEN', 'CH4', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'NOx', 'O_3', 'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4127 entries, 87457 to 158286
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   date        4127 non-null   object
1   BEN         4127 non-null   float64
2   CH4         4127 non-null   float64
3   CO          4127 non-null   float64
4   EBE         4127 non-null   float64
5   NMHC        4127 non-null   float64
6   NO          4127 non-null   float64
7   NO_2        4127 non-null   float64
8   NOx         4127 non-null   float64
9   O_3         4127 non-null   float64
10  PM10        4127 non-null   float64
11  PM25        4127 non-null   float64
12  SO_2        4127 non-null   float64
13  TCH         4127 non-null   float64
14  TOL         4127 non-null   float64
15  station     4127 non-null   int64
dtypes: float64(14), int64(1), object(1)
memory usage: 548.1+ KB
```

In [6]: `data=df[['CO', 'station']]`
`data`

Out[6]:


	CO	station
87457	0.3	28079008
87462	0.2	28079024
87481	0.2	28079008
87486	0.2	28079024
87505	0.2	28079008
...
158238	0.2	28079024
158257	0.3	28079008
158262	0.2	28079024
158281	0.2	28079008
158286	0.2	28079024

4127 rows × 2 columns

Line chart

In [7]: `data.plot.line(subplots=True)`


Out[7]: `array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)`



Line chart

In [8]: `data.plot.line()`

Out[8]: `<AxesSubplot:>`




Bar chart

In [9]: `b=data[0:50]`

In [10]: `b.plot.bar()`


Out[10]: `<AxesSubplot:>`



Histogram

```
In [11]: data.plot.hist()
```


```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

```
In [12]: data.plot.area()
```


```
Out[12]: <AxesSubplot:>
```



Box chart

```
In [13]: data.plot.box()
```


```
Out[13]: <AxesSubplot:>
```



Pie chart

```
In [14]: b.plot.pie(y='station' )
```


```
Out[14]: <AxesSubplot:ylabel='station'>
```



Scatter chart

```
In [15]: data.plot.scatter(x='CO' ,y='station')
```

Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>



In [16]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4127 entries, 87457 to 158286
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        4127 non-null   object
1   BEN         4127 non-null   float64
2   CH4         4127 non-null   float64
3   CO          4127 non-null   float64
4   EBE         4127 non-null   float64
5   NMHC        4127 non-null   float64
6   NO          4127 non-null   float64
7   NO_2        4127 non-null   float64
8   NOx         4127 non-null   float64
9   O_3         4127 non-null   float64
10  PM10        4127 non-null   float64
11  PM25        4127 non-null   float64
12  SO_2        4127 non-null   float64
13  TCH         4127 non-null   float64
14  TOL         4127 non-null   float64
15  station     4127 non-null   int64
dtypes: float64(14), int64(1), object(1)
memory usage: 548.1+ KB
```

In [17]: `df.columns`

Out[17]: Index(['date', 'BEN', 'CH4', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'NOx', 'O_3', 'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')

In [18]: `df.describe()`

	BEN	CH4	CO	EBE	NMHC	NO	NO_2	
count	4127.000000	4127.000000	4127.000000	4127.000000	4127.000000	4127.000000	4127.000000	4
mean	0.919918	1.323732	0.417858	0.578168	0.097269	41.785316	58.069057	
std	1.123078	0.215742	0.342871	0.962000	0.094035	71.118499	38.974112	
min	0.100000	1.100000	0.100000	0.100000	0.000000	1.000000	1.000000	
25%	0.300000	1.180000	0.200000	0.100000	0.050000	3.000000	30.000000	

	BEN	CH4	CO	EBE	NMHC	NO	NO_2	
50%	0.600000	1.270000	0.300000	0.300000	0.080000	16.000000	54.000000	
75%	1.100000	1.400000	0.500000	0.700000	0.110000	50.000000	78.000000	
max	19.600000	3.630000	4.900000	16.700001	1.420000	879.000000	349.000000	1

In [19]:

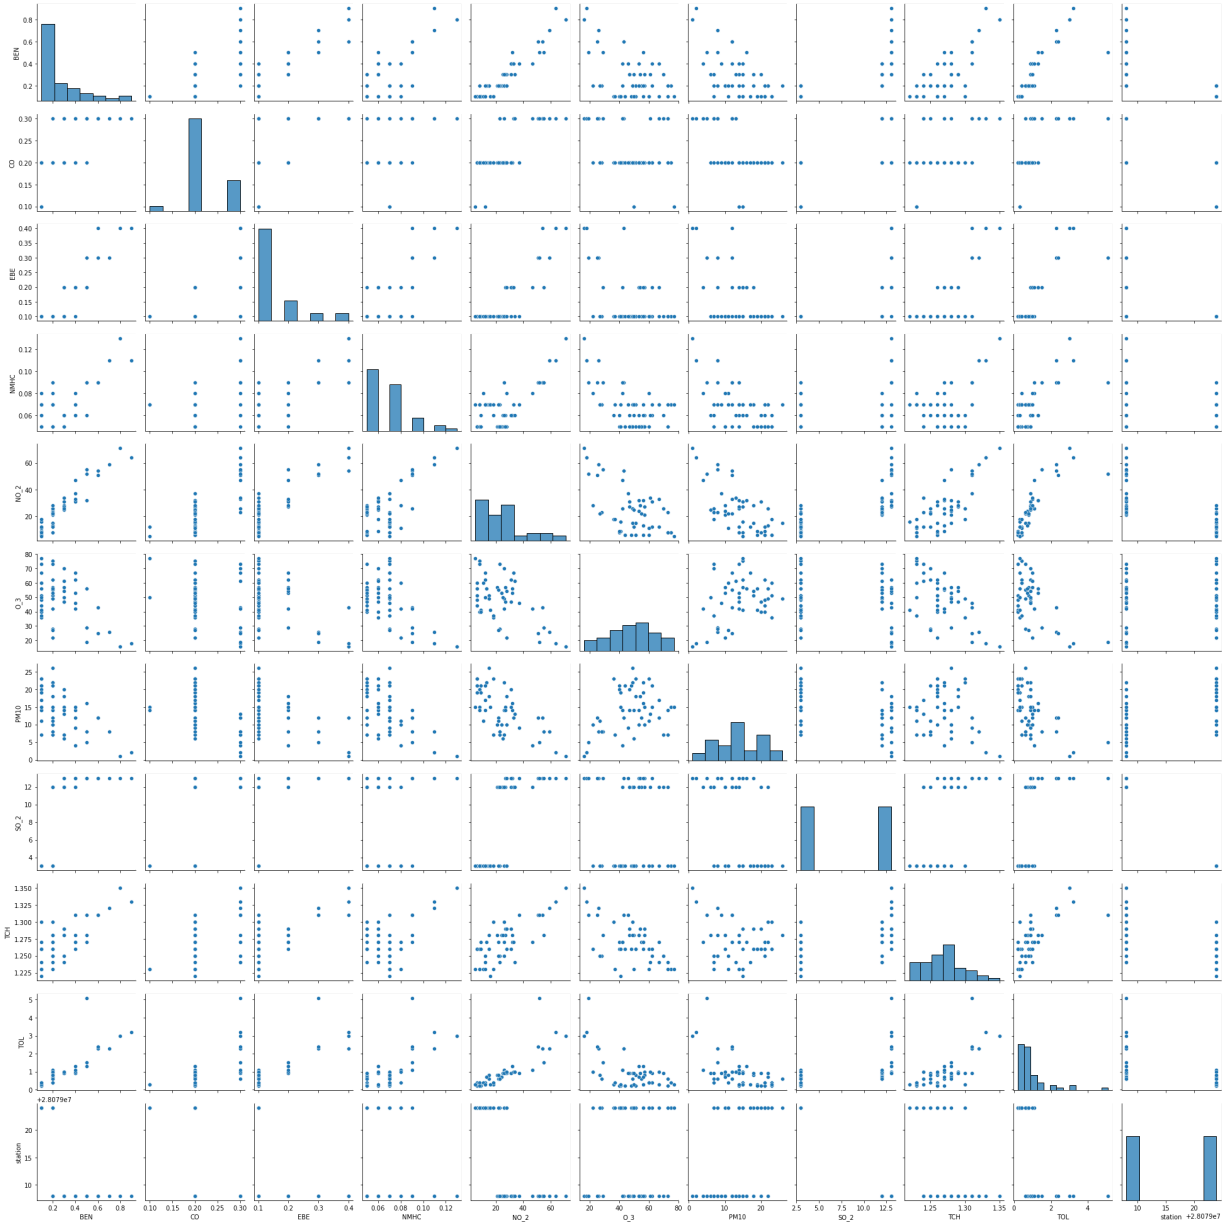
```
df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2','O_3',  
        'PM10','SO_2', 'TCH', 'TOL', 'station']]
```

EDA AND VISUALIZATION

In [20]:

```
sns.pairplot(df1[0:50])
```

Out[20]: <seaborn.axisgrid.PairGrid at 0x2eda3b5af10>




In [21]:

```
sns.distplot(df1['station'])
```



C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
 warnings.warn(msg, FutureWarning)

Out[21]: <AxesSubplot:xlabel='station', ylabel='Density'>



In [22]: `sns.heatmap(df1.corr())`

Out[22]: <AxesSubplot:>



TO TRAIN THE MODEL AND MODEL BUILDING

In [23]: `x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
 'PM10', 'SO_2', 'TCH', 'TOL']]
 y=df['station']`

In [24]: `from sklearn.model_selection import train_test_split
 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)`

Linear Regression

```
In [25]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[25]: LinearRegression()

```
In [26]: lr.intercept_
```

Out[26]: 28079039.742447674


```
In [27]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[27]:

	Co-efficient
BEN	1.085227
CO	-2.089988
EBE	-2.720277
NMHC	36.240965
NO_2	-0.163846
O_3	-0.083555
PM10	0.353406
SO_2	-0.250277
TCH	-13.987635
TOL	0.231713

```
In [28]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[28]: <matplotlib.collections.PathCollection at 0x2edac169580>



ACCURACY

```
In [29]: lr.score(x_test,y_test)
```

```
Out[29]: 0.5900771355947994
```

```
In [30]: lr.score(x_train,y_train)
```

```
Out[30]: 0.6106445957418349
```

Ridge and Lasso

```
In [31]: from sklearn.linear_model import Ridge,Lasso
```

```
In [32]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
Out[32]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [33]: rr.score(x_test,y_test)
```

```
Out[33]: 0.5715248038847451
```

```
In [34]: rr.score(x_train,y_train)
```

```
Out[34]: 0.5938530282276366
```

```
In [35]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[35]: Lasso(alpha=10)
```

```
In [36]: la.score(x_train,y_train)
```

```
Out[36]: 0.40482712909937557
```

Accuracy(Lasso)

```
In [37]: la.score(x_test,y_test)
```

```
Out[37]: 0.39778549225268334
```

Elastic Net

```
In [38]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(x_train,y_train)
```

Out[38]: ElasticNet()

```
In [39]: en.coef_
```

Out[39]: array([-0. , -0. , -0. , 0. , -0.16620446,
 -0.0579603 , 0.33131595, -0.34871237, -0. , 0.11185378])

```
In [40]: en.intercept_
```

Out[40]: 28079022.7783964

```
In [41]: prediction=en.predict(x_test)
```

```
In [42]: en.score(x_test,y_test)
```

Out[42]: 0.4418261548618335

Evaluation Metrics

```
In [43]: from sklearn import metrics
         print(metrics.mean_absolute_error(y_test,prediction))
         print(metrics.mean_squared_error(y_test,prediction))
         print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

5.145347586020383
35.72198583137868
5.976787249967885

Logistic Regression

```
In [44]: from sklearn.linear_model import LogisticRegression
```

```
In [45]: feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2','O_3',  
                          'PM10','SO_2', 'TCH', 'TOL']]
         target_vector=df['station']
```

```
In [46]: feature_matrix.shape
```

Out[46]: (4127, 10)

```
In [47]: target_vector.shape
```

Out[47]: (4127,)

```
In [48]: from sklearn.preprocessing import StandardScaler
```

```
In [49]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [50]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[50]: LogisticRegression(max_iter=10000)
```

```
In [51]: observation=[[1,2,3,4,5,6,7,8,9,10]]
```

```
In [52]: prediction=logr.predict(observation)
print(prediction)
```

```
[28079008]
```

```
In [53]: logr.classes_
```

```
Out[53]: array([28079008, 28079024], dtype=int64)
```

```
In [54]: logr.score(fs,target_vector)
```

```
Out[54]: 0.9437848315968016
```

```
In [55]: logr.predict_proba(observation)[0][0]
```

```
Out[55]: 0.999999999725541
```

```
In [56]: logr.predict_proba(observation)
```

```
Out[56]: array([[1.00000000e+00, 2.74458959e-11]])
```

Random Forest

```
In [57]: from sklearn.ensemble import RandomForestClassifier
```

```
In [58]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[58]: RandomForestClassifier()
```

```
In [59]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]
}
```

```

In [60]: from sklearn.model_selection import GridSearchCV
         grid_search = GridSearchCV(estimator=rfc, param_grid=parameters, cv=2, scoring="accuracy")
         grid_search.fit(x_train, y_train)

Out[60]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')

In [61]: grid_search.best_score_

Out[61]: 0.9698753462603877

In [62]: rfc_best = grid_search.best_estimator_

In [63]: from sklearn.tree import plot_tree

         plt.figure(figsize=(80,40))
         plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b', 'c'],

Out[63]: [Text(1715.85, 1993.2, 'CO <= 0.25\ngini = 0.499\nsamples = 1840\nvalue = [1509, 137
9]\nclass = a'),
         Text(725.4, 1630.8000000000002, 'SO_2 <= 5.5\ngini = 0.305\nsamples = 576\nvalue =
[165, 715]\nclass = b'),
         Text(334.79999999999995, 1268.4, 'SO_2 <= 2.5\ngini = 0.193\nsamples = 525\nvalue =
[86, 709]\nclass = b'),
         Text(223.2, 906.0, 'gini = 0.0\nsamples = 26\nvalue = [42, 0]\nclass = a'),
         Text(446.4, 906.0, 'EBE <= 0.15\ngini = 0.11\nsamples = 499\nvalue = [44, 709]\nclass = b'),
         Text(223.2, 543.5999999999999, 'NO_2 <= 24.5\ngini = 0.05\nsamples = 441\nvalue =
[17, 640]\nclass = b'),
         Text(111.6, 181.19999999999982, 'gini = 0.012\nsamples = 325\nvalue = [3, 483]\nclass = b'),
         Text(334.79999999999995, 181.19999999999982, 'gini = 0.15\nsamples = 116\nvalue =
[14, 157]\nclass = b'),
         Text(669.5999999999999, 543.5999999999999, 'PM10 <= 11.5\ngini = 0.404\nsamples = 5
8\nvalue = [27, 69]\nclass = b'),
         Text(558.0, 181.19999999999982, 'gini = 0.287\nsamples = 13\nvalue = [19, 4]\nclass = a'),
         Text(781.1999999999999, 181.19999999999982, 'gini = 0.195\nsamples = 45\nvalue =
[8, 65]\nclass = b'),
         Text(1116.0, 1268.4, 'PM10 <= 15.5\ngini = 0.131\nsamples = 51\nvalue = [79, 6]\nclass = a'),
         Text(1004.4, 906.0, 'NMHC <= 0.025\ngini = 0.074\nsamples = 46\nvalue = [75, 3]\nclass = a'),
         Text(892.8, 543.5999999999999, 'gini = 0.444\nsamples = 7\nvalue = [6, 3]\nclass = a'),
         Text(1116.0, 543.5999999999999, 'gini = 0.0\nsamples = 39\nvalue = [69, 0]\nclass = a'),
         Text(1227.6, 906.0, 'gini = 0.49\nsamples = 5\nvalue = [4, 3]\nclass = a'),
         Text(2706.2999999999997, 1630.8000000000002, 'EBE <= 0.15\ngini = 0.443\nsamples =
1264\nvalue = [1344, 664]\nclass = a'),
         Text(1841.3999999999999, 1268.4, 'O_3 <= 14.5\ngini = 0.392\nsamples = 184\nvalue =
[80, 219]\nclass = b'),
         Text(1450.8, 906.0, 'TOL <= 0.85\ngini = 0.028\nsamples = 91\nvalue = [2, 139]\nclass = b'),
         Text(1339.1999999999998, 543.5999999999999, 'gini = 0.219\nsamples = 5\nvalue = [1,
7]\nclass = b'),
         Text(1562.3999999999999, 543.5999999999999, 'PM10 <= 6.5\ngini = 0.015\nsamples = 8
6\nvalue = [1, 132]\nclass = b'),
         Text(1450.8, 181.19999999999982, 'gini = 0.278\nsamples = 5\nvalue = [1, 5]\nclass

```

[illegible]

Conclusion

Scores

Linear Regression

```
In [64]: lr.score(x_test,y_test)
```

```
Out[64]: 0.5900771355947994
```

```
In [65]: lr.score(x_train,y_train)
```

```
Out[65]: 0.6106445957418349
```

Lasso

```
In [66]: la.score(x_test,y_test)
```

```
Out[66]: 0.39778549225268334
```

Ridge

```
In [67]: rr.score(x_test,y_test)
```

```
Out[67]: 0.5715248038847451
```

```
In [68]: rr.score(x_train,y_train)
```

```
Out[68]: 0.5938530282276366
```

Elastic Net

```
In [69]: en.score(x_test,y_test)
```

```
Out[69]: 0.4418261548618335
```

Logistic Regression

```
In [70]: logr.score(fs,target_vector)
```

```
Out[70]: 0.9437848315968016
```


Random Forest

In [71]:

```
grid_search.best_score_
```

Out[71]: 0.9698753462603877

From the above data, we can conclude that random forest regression is preferable to other regression types

In []: