

Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

In [2]:

```
df=pd.read_csv("i2015.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	sta
0	2015-10-01 01:00:00	NaN	0.8	NaN	NaN	90.0	82.0	NaN	NaN	NaN	10.0	NaN	NaN	2807
1	2015-10-01 01:00:00	2.0	0.8	1.6	0.33	40.0	95.0	4.0	37.0	24.0	12.0	1.83	8.3	2807
2	2015-10-01 01:00:00	3.1	NaN	1.8	NaN	29.0	97.0	NaN	NaN	NaN	NaN	NaN	7.1	2807
3	2015-10-01 01:00:00	NaN	0.6	NaN	NaN	30.0	103.0	2.0	NaN	NaN	NaN	NaN	NaN	2807
4	2015-10-01 01:00:00	NaN	NaN	NaN	NaN	95.0	96.0	2.0	NaN	NaN	9.0	NaN	NaN	2807
...
210091	2015-08-01 00:00:00	NaN	0.2	NaN	NaN	11.0	33.0	53.0	NaN	NaN	NaN	NaN	NaN	2807
210092	2015-08-01 00:00:00	NaN	0.2	NaN	NaN	1.0	5.0	NaN	26.0	NaN	10.0	NaN	NaN	2807
210093	2015-08-01 00:00:00	NaN	NaN	NaN	NaN	1.0	7.0	74.0	NaN	NaN	NaN	NaN	NaN	2807
210094	2015-08-01 00:00:00	NaN	NaN	NaN	NaN	3.0	7.0	65.0	NaN	NaN	NaN	NaN	NaN	2807
210095	2015-08-01 00:00:00	NaN	NaN	NaN	NaN	1.0	9.0	54.0	29.0	NaN	NaN	NaN	NaN	2807
210096 rows × 14 columns														

Data Cleaning and Data Preprocessing

In [3]:

```
df=df.dropna()
```

In [4]:

```
df.columns
```

Out[4]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
      'SO_2', 'TCH', 'TOL', 'station'],  
      dtype='object')
```

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 16026 entries, 1 to 210078  
Data columns (total 14 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   date        16026 non-null   object  
1   BEN         16026 non-null   float64  
2   CO          16026 non-null   float64  
3   EBE         16026 non-null   float64  
4   NMHC        16026 non-null   float64  
5   NO          16026 non-null   float64  
6   NO_2        16026 non-null   float64  
7   O_3         16026 non-null   float64  
8   PM10        16026 non-null   float64  
9   PM25        16026 non-null   float64  
10  SO_2        16026 non-null   float64  
11  TCH         16026 non-null   float64  
12  TOL         16026 non-null   float64  
13  station     16026 non-null   int64  
dtypes: float64(12), int64(1), object(1)  
memory usage: 1.8+ MB
```

In [6]:

```
data=df[['CO' , 'station']]  
data
```

Out[6]:


	CO	station
1	0.8	28079008
6	0.3	28079024
25	0.7	28079008
30	0.3	28079024
49	0.8	28079008
...
210030	0.1	28079024
210049	0.3	28079008
210054	0.1	28079024
210073	0.3	28079008
210078	0.1	28079024

16026 rows × 2 columns

Line chart

```
In [7]: data.plot.line(subplots=True)
```


```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



Line chart

```
In [8]: data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```




Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```


```
Out[10]: <AxesSubplot:>
```



Histogram

```
In [11]: data.plot.hist()
```


```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

```
In [12]: data.plot.area()
```


```
Out[12]: <AxesSubplot:>
```



Box chart

In [13]: `data.plot.box()`


Out[13]: `<AxesSubplot:>`



Pie chart

In [14]: `b.plot.pie(y='station')`


Out[14]: `<AxesSubplot:ylabel='station'>`



Scatter chart

```
In [15]: data.plot.scatter(x='CO' ,y='station')
```

Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>



```
In [16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 16026 entries, 1 to 210078
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        16026 non-null  object
1   BEN         16026 non-null  float64
2   CO          16026 non-null  float64
3   EBE         16026 non-null  float64
4   NMHC        16026 non-null  float64
5   NO          16026 non-null  float64
6   NO_2        16026 non-null  float64
7   O_3         16026 non-null  float64
8   PM10        16026 non-null  float64
9   PM25        16026 non-null  float64
10  SO_2        16026 non-null  float64
11  TCH         16026 non-null  float64
12  TOL         16026 non-null  float64
13  station     16026 non-null  int64
dtypes: float64(12), int64(1), object(1)
memory usage: 1.8+ MB
```

```
In [17]: df.columns
```

```
Out[17]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
               'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

```
In [18]: df.describe()
```

Out[18]:

	BEN	CO	EBE	NMHC	NO	NO_2	
count	16026.000000	16026.000000	16026.000000	16026.000000	16026.000000	16026.000000	16026.00
mean	0.504823	0.380594	0.394247	0.123099	23.842256	40.948771	48.08
std	0.716896	0.260805	0.678592	0.092368	51.255660	33.236098	35.84
min	0.100000	0.100000	0.100000	0.000000	1.000000	1.000000	1.00
25%	0.100000	0.200000	0.100000	0.070000	1.000000	14.000000	15.00
50%	0.200000	0.300000	0.100000	0.100000	6.000000	35.000000	46.00
75%	0.700000	0.500000	0.400000	0.140000	24.000000	60.000000	73.00

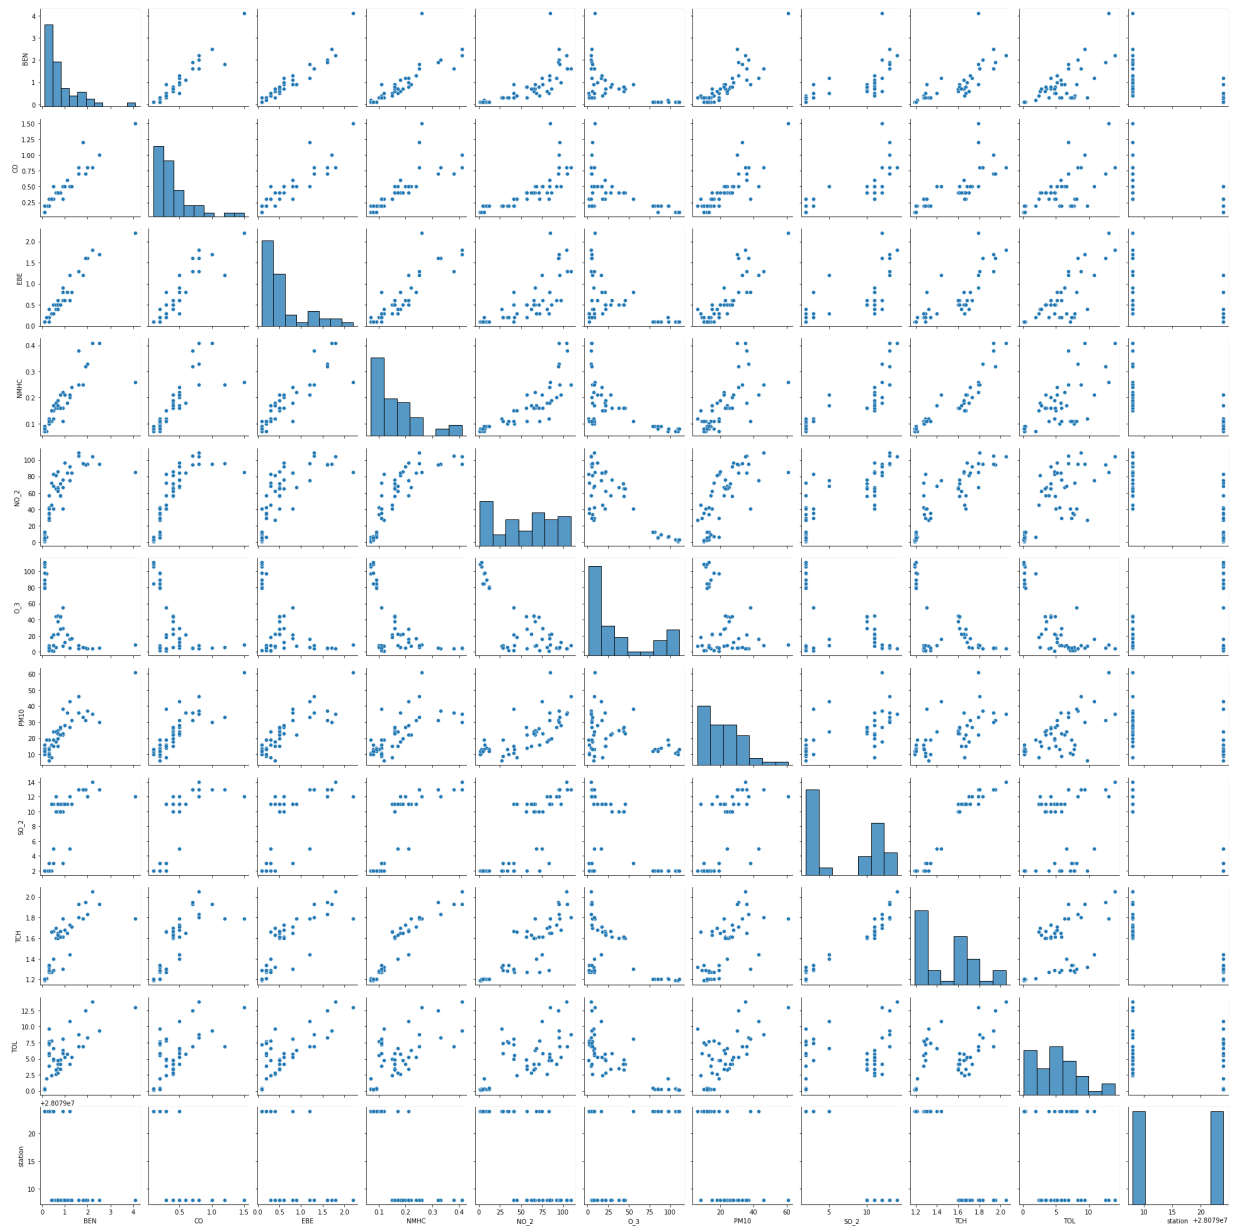
	BEN	CO	EBE	NMHC	NO	NO_2	NO_2
max	17.700001	4.500000	12.100000	1.090000	960.000000	369.000000	217.00

```
In [19]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2','O_3',
              'PM10','SO_2', 'TCH', 'TOL', 'station']]
```

EDA AND VISUALIZATION

```
In [20]: sns.pairplot(df1[0:50])
```

```
Out[20]: <seaborn.axisgrid.PairGrid at 0x2b88a8dc880>
```




```
In [21]: sns.distplot(df1['station'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar f


lexibility) or `histplot` (an axes-level function for histograms).
 warnings.warn(msg, FutureWarning)

Out[21]: <AxesSubplot:xlabel='station', ylabel='Density'>



In [22]: `sns.heatmap(df1.corr())`

Out[22]: <AxesSubplot:>



TO TRAIN THE MODEL AND MODEL BUILDING

In [23]: `x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
 'PM10', 'SO_2', 'TCH', 'TOL']]
 y=df['station']`

In [24]: `from sklearn.model_selection import train_test_split
 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)`

Linear Regression

```
In [25]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[25]: LinearRegression()

```
In [26]: lr.intercept_
```

Out[26]: 28079037.971513327


```
In [27]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[27]:

	Co-efficient
BEN	4.829039
CO	-7.201094
EBE	-1.221655
NMHC	24.353949
NO_2	-0.001239
O_3	-0.017437
PM10	0.057771
SO_2	-1.173156
TCH	-10.943129
TOL	-0.070548

```
In [28]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[28]: <matplotlib.collections.PathCollection at 0x2b892efa7c0>



ACCURACY

```
In [29]: lr.score(x_test,y_test)
```

```
Out[29]: 0.8058280829064649
```

```
In [30]: lr.score(x_train,y_train)
```

```
Out[30]: 0.8089532029954628
```

Ridge and Lasso

```
In [31]: from sklearn.linear_model import Ridge,Lasso
```

```
In [32]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
Out[32]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [33]: rr.score(x_test,y_test)
```

```
Out[33]: 0.8042402331444318
```

```
In [34]: rr.score(x_train,y_train)
```

```
Out[34]: 0.8068821169448739
```

```
In [35]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[35]: Lasso(alpha=10)
```

```
In [36]: la.score(x_train,y_train)
```

```
Out[36]: 0.6361709027833888
```

Accuracy(Lasso)

```
In [37]: la.score(x_test,y_test)
```

```
Out[37]: 0.6343603719349661
```

Elastic Net

```
In [38]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(x_train,y_train)
```

Out[38]: ElasticNet()

```
In [39]: en.coef_
```

Out[39]: array([0. , -0. , -0. , 0. , -0. ,
 -0.01573327, 0.07181082, -1.25810757, -0. , 0.15780668])

```
In [40]: en.intercept_
```

Out[40]: 28079024.1919828

```
In [41]: prediction=en.predict(x_test)
```

```
In [42]: en.score(x_test,y_test)
```

Out[42]: 0.7348213384433098

Evaluation Metrics

```
In [43]: from sklearn import metrics
         print(metrics.mean_absolute_error(y_test,prediction))
         print(metrics.mean_squared_error(y_test,prediction))
         print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

3.234674387959339
16.967628463726648
4.11917812964269

Logistic Regression

```
In [44]: from sklearn.linear_model import LogisticRegression
```

```
In [45]: feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2','O_3',  
                          'PM10','SO_2', 'TCH', 'TOL']]
         target_vector=df['station']
```

```
In [46]: feature_matrix.shape
```

Out[46]: (16026, 10)

```
In [47]: target_vector.shape
```

Out[47]: (16026,)

```
In [48]: from sklearn.preprocessing import StandardScaler
```

```
In [49]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [50]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[50]: LogisticRegression(max_iter=10000)
```

```
In [51]: observation=[[1,2,3,4,5,6,7,8,9,10]]
```

```
In [52]: prediction=logr.predict(observation)
print(prediction)
```

```
[28079008]
```

```
In [53]: logr.classes_
```

```
Out[53]: array([28079008, 28079024], dtype=int64)
```

```
In [54]: logr.score(fs,target_vector)
```

```
Out[54]: 0.9947585174092101
```

```
In [55]: logr.predict_proba(observation)[0][0]
```

```
Out[55]: 1.0
```

```
In [56]: logr.predict_proba(observation)
```

```
Out[56]: array([[1.00000000e+00, 5.69793111e-39]])
```

Random Forest

```
In [57]: from sklearn.ensemble import RandomForestClassifier
```

```
In [58]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[58]: RandomForestClassifier()
```

```
In [59]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]
}
```

```

In [60]: from sklearn.model_selection import GridSearchCV
         grid_search = GridSearchCV(estimator=rfc, param_grid=parameters, cv=2, scoring="accuracy")
         grid_search.fit(x_train, y_train)

Out[60]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')

In [61]: grid_search.best_score_

Out[61]: 0.9948297379211981

In [62]: rfc_best = grid_search.best_estimator_

In [63]: from sklearn.tree import plot_tree

         plt.figure(figsize=(80,40))
         plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b', 'c'],

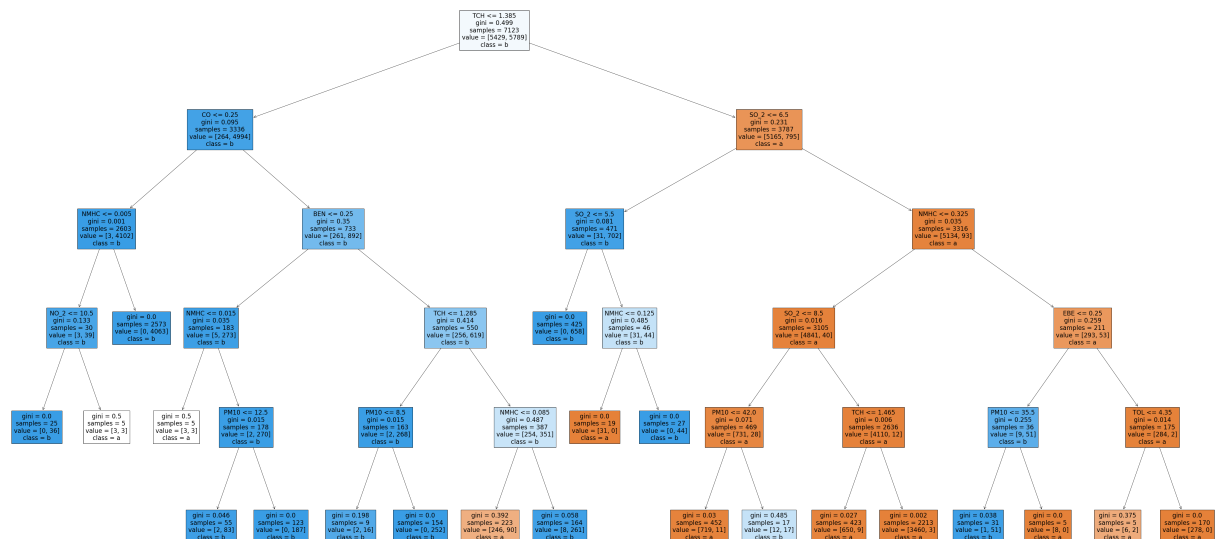
Out[63]: [Text(1801.5428571428572, 1993.2, 'TCH <= 1.385\ngini = 0.499\nsamples = 7123\nvalue
= [5429, 5789]\nclass = b'),
         Text(797.1428571428571, 1630.8000000000002, 'CO <= 0.25\ngini = 0.095\nsamples = 33
36\nvalue = [264, 4994]\nclass = b'),
         Text(382.62857142857143, 1268.4, 'NMHC <= 0.005\ngini = 0.001\nsamples = 2603\nvalu
e = [3, 4102]\nclass = b'),
         Text(255.0857142857143, 906.0, 'NO_2 <= 10.5\ngini = 0.133\nsamples = 30\nvalue =
[3, 39]\nclass = b'),
         Text(127.54285714285714, 543.5999999999999, 'gini = 0.0\nsamples = 25\nvalue = [0,
36]\nclass = b'),
         Text(382.62857142857143, 543.5999999999999, 'gini = 0.5\nsamples = 5\nvalue = [3,
3]\nclass = a'),
         Text(510.1714285714286, 906.0, 'gini = 0.0\nsamples = 2573\nvalue = [0, 4063]\nclass
= b'),
         Text(1211.6571428571428, 1268.4, 'BEN <= 0.25\ngini = 0.35\nsamples = 733\nvalue =
[261, 892]\nclass = b'),
         Text(765.2571428571429, 906.0, 'NMHC <= 0.015\ngini = 0.035\nsamples = 183\nvalue =
[5, 273]\nclass = b'),
         Text(637.7142857142858, 543.5999999999999, 'gini = 0.5\nsamples = 5\nvalue = [3, 3]
\nclass = a'),
         Text(892.8, 543.5999999999999, 'PM10 <= 12.5\ngini = 0.015\nsamples = 178\nvalue =
[2, 270]\nclass = b'),
         Text(765.2571428571429, 181.19999999999982, 'gini = 0.046\nsamples = 55\nvalue =
[2, 83]\nclass = b'),
         Text(1020.3428571428572, 181.19999999999982, 'gini = 0.0\nsamples = 123\nvalue =
[0, 187]\nclass = b'),
         Text(1658.057142857143, 906.0, 'TCH <= 1.285\ngini = 0.414\nsamples = 550\nvalue =
[256, 619]\nclass = b'),
         Text(1402.9714285714285, 543.5999999999999, 'PM10 <= 8.5\ngini = 0.015\nsamples = 1
63\nvalue = [2, 268]\nclass = b'),
         Text(1275.4285714285716, 181.19999999999982, 'gini = 0.198\nsamples = 9\nvalue =
[2, 16]\nclass = b'),
         Text(1530.5142857142857, 181.19999999999982, 'gini = 0.0\nsamples = 154\nvalue =
[0, 252]\nclass = b'),
         Text(1913.142857142857, 543.5999999999999, 'NMHC <= 0.085\ngini = 0.487\nsamples =
387\nvalue = [254, 351]\nclass = b'),
         Text(1785.6, 181.19999999999982, 'gini = 0.392\nsamples = 223\nvalue = [246, 90]\nc
lass = a'),
         Text(2040.6857142857143, 181.19999999999982, 'gini = 0.058\nsamples = 164\nvalue =
[8, 261]\nclass = b'),
         Text(2805.942857142857, 1630.8000000000002, 'SO_2 <= 6.5\ngini = 0.231\nsamples = 3

```

```

787\nvalue = [5165, 795]\nclasse = a'),
Text(2168.2285714285713, 1268.4, 'SO_2 <= 5.5\nngini = 0.081\nnsamples = 471\nvalue = [31, 702]\nclasse = b'),
Text(2040.6857142857143, 906.0, 'gini = 0.0\nnsamples = 425\nvalue = [0, 658]\nclasse = b'),
Text(2295.7714285714287, 906.0, 'NMHC <= 0.125\nngini = 0.485\nnsamples = 46\nvalue = [31, 44]\nclasse = b'),
Text(2168.2285714285713, 543.5999999999999, 'gini = 0.0\nnsamples = 19\nvalue = [31, 0]\nclasse = a'),
Text(2423.3142857142857, 543.5999999999999, 'gini = 0.0\nnsamples = 27\nvalue = [0, 44]\nclasse = b'),
Text(3443.657142857143, 1268.4, 'NMHC <= 0.325\nngini = 0.035\nnsamples = 3316\nvalue = [5134, 93]\nclasse = a'),
Text(2933.4857142857145, 906.0, 'SO_2 <= 8.5\nngini = 0.016\nnsamples = 3105\nvalue = [4841, 40]\nclasse = a'),
Text(2678.4, 543.5999999999999, 'PM10 <= 42.0\nngini = 0.071\nnsamples = 469\nvalue = [731, 28]\nclasse = a'),
Text(2550.857142857143, 181.19999999999998, 'gini = 0.03\nnsamples = 452\nvalue = [7 19, 11]\nclasse = a'),
Text(2805.942857142857, 181.19999999999998, 'gini = 0.485\nnsamples = 17\nvalue = [1 2, 17]\nclasse = b'),
Text(3188.5714285714284, 543.5999999999999, 'TCH <= 1.465\nngini = 0.006\nnsamples = 2636\nvalue = [4110, 12]\nclasse = a'),
Text(3061.0285714285715, 181.19999999999998, 'gini = 0.027\nnsamples = 423\nvalue = [650, 9]\nclasse = a'),
Text(3316.114285714286, 181.19999999999998, 'gini = 0.002\nnsamples = 2213\nvalue = [3460, 3]\nclasse = a'),
Text(3953.8285714285716, 906.0, 'EBE <= 0.25\nngini = 0.259\nnsamples = 211\nvalue = [293, 53]\nclasse = a'),
Text(3698.7428571428572, 543.5999999999999, 'PM10 <= 35.5\nngini = 0.255\nnsamples = 36\nvalue = [9, 51]\nclasse = b'),
Text(3571.2, 181.19999999999998, 'gini = 0.038\nnsamples = 31\nvalue = [1, 51]\nclasse = b'),
Text(3826.285714285714, 181.19999999999998, 'gini = 0.0\nnsamples = 5\nvalue = [8, 0]\nclasse = a'),
Text(4208.914285714286, 543.5999999999999, 'TOL <= 4.35\nngini = 0.014\nnsamples = 17 5\nvalue = [284, 2]\nclasse = a'),
Text(4081.3714285714286, 181.19999999999998, 'gini = 0.375\nnsamples = 5\nvalue = [6, 2]\nclasse = a'),
Text(4336.457142857143, 181.19999999999998, 'gini = 0.0\nnsamples = 170\nvalue = [27 8, 0]\nclasse = a')]

```



Conclusion

Scores

Linear Regression

```
In [64]: lr.score(x_test,y_test)
```

```
Out[64]: 0.8058280829064649
```

```
In [65]: lr.score(x_train,y_train)
```

```
Out[65]: 0.8089532029954628
```

Lasso

```
In [66]: la.score(x_test,y_test)
```

```
Out[66]: 0.6343603719349661
```

Ridge

```
In [67]: rr.score(x_test,y_test)
```

```
Out[67]: 0.8042402331444318
```

```
In [68]: rr.score(x_train,y_train)
```

```
Out[68]: 0.8068821169448739
```

Elastic Net

```
In [69]: en.score(x_test,y_test)
```

```
Out[69]: 0.7348213384433098
```

Logistic Regression

```
In [70]: logr.score(fs,target_vector)
```

```
Out[70]: 0.9947585174092101
```

Random Forest

```
In [71]: grid_search.best_score_
```


Out[71]: 0.9948297379211981

From the above data, we can conclude that random forest regression and logistic regression is preferable to other regression types

In []: