

# Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

# Importing Datasets

```
In [2]: df=pd.read_csv("stations.csv")
df
```

Out[2]:

|    | id       | name                 | address                                          | lon       | lat       | elevation |
|----|----------|----------------------|--------------------------------------------------|-----------|-----------|-----------|
| 0  | 28079004 | Pza. de España       | Plaza de España                                  | -3.712247 | 40.423853 | 635       |
| 1  | 28079008 | Escuelas Aguirre     | Entre C/ Alcalá y C/ O' Donell                   | -3.682319 | 40.421564 | 670       |
| 2  | 28079011 | Avda. Ramón y Cajal  | Avda. Ramón y Cajal esq. C/ Príncipe de Vergara  | -3.677356 | 40.451475 | 708       |
| 3  | 28079016 | Arturo Soria         | C/ Arturo Soria esq. C/ Vizconde de los Asilos   | -3.639233 | 40.440047 | 693       |
| 4  | 28079017 | Villaverde           | C/. Juan Peñalver                                | -3.713322 | 40.347139 | 604       |
| 5  | 28079018 | Farolillo            | Calle Farolillo - C/Ervigio                      | -3.731853 | 40.394781 | 630       |
| 6  | 28079024 | Casa de Campo        | Casa de Campo (Terminal del Teleférico)          | -3.747347 | 40.419356 | 642       |
| 7  | 28079027 | Barajas Pueblo       | C/. Júpiter, 21 (Barajas)                        | -3.580031 | 40.476928 | 621       |
| 8  | 28079035 | Pza. del Carmen      | Plaza del Carmen esq. Tres Cruces.               | -3.703172 | 40.419208 | 659       |
| 9  | 28079036 | Moratalaz            | Avd. Moratalaz esq. Camino de los Vinateros      | -3.645306 | 40.407947 | 685       |
| 10 | 28079038 | Cuatro Caminos       | Avda. Pablo Iglesias esq. C/ Marqués de Lema     | -3.707128 | 40.445544 | 698       |
| 11 | 28079039 | Barrio del Pilar     | Avd. Betanzos esq. C/ Monforte de Lemos          | -3.711542 | 40.478228 | 674       |
| 12 | 28079040 | Vallecas             | C/ Arroyo del Olivar esq. C/ Río Grande.         | -3.651522 | 40.388153 | 677       |
| 13 | 28079047 | Mendez Alvaro        | C/ Juan de Mariana / Pza. Amanecer Mendez Alvaro | -3.686825 | 40.398114 | 599       |
| 14 | 28079048 | Castellana           | C/ Jose Gutierrez Abascal                        | -3.690367 | 40.439897 | 676       |
| 15 | 28079049 | Parque del Retiro    | Paseo Venezuela- Casa de Vacas                   | -3.682583 | 40.414444 | 662       |
| 16 | 28079050 | Plaza Castilla       | Plaza Castilla (Canal)                           | -3.688769 | 40.465572 | 728       |
| 17 | 28079054 | Ensanche de Vallecas | Avda La Gavia / Avda. Las Suertes                | -3.612117 | 40.372933 | 627       |
| 18 | 28079055 | Urb. Embajada        | C/ Riaño (Barajas)                               | -3.580747 | 40.462531 | 618       |
| 19 | 28079056 | Pza. Fernández       | Pza. Fernández Ladreda - Avda.                   | -3.718728 | 40.384964 | 604       |

|    | id       | name          | address                                           | lon       | lat       | elevation |
|----|----------|---------------|---------------------------------------------------|-----------|-----------|-----------|
|    |          | Ladreda       | Oporto                                            |           |           |           |
| 20 | 28079057 | Sanchinarro   | C/ Princesa de Eboli esq C/ Maria Tudor           | -3.660503 | 40.494208 | 700       |
| 21 | 28079058 | El Pardo      | Avda. La Guardia                                  | -3.774611 | 40.518058 | 615       |
| 22 | 28079059 | Juan Carlos I | Parque Juan Carlos I (frente oficinas mantenim... | -3.609072 | 40.465250 | 660       |
| 23 | 28079060 | Tres Olivos   | Plaza Tres Olivos                                 | -3.689761 | 40.500589 | 715       |

## Data Cleaning and Data Preprocessing

In [3]:

df=df.dropna()

In [8]:

df.columns

Out[8]: Index(['id', 'name', 'address', 'lon', 'lat', 'elevation'], dtype='object')

In [9]:

df.info()  
  
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 24 entries, 0 to 23  
Data columns (total 6 columns):  
# Column Non-Null Count Dtype  
--- ---  
0 id 24 non-null int64  
1 name 24 non-null object  
2 address 24 non-null object  
3 lon 24 non-null float64  
4 lat 24 non-null float64  
5 elevation 24 non-null int64  
dtypes: float64(2), int64(2), object(2)  
memory usage: 1.3+ KB

In [10]:

data=df[['id' , 'lon']]  
data

Out[10]:


|   | id       | lon       |
|---|----------|-----------|
| 0 | 28079004 | -3.712247 |
| 1 | 28079008 | -3.682319 |
| 2 | 28079011 | -3.677356 |
| 3 | 28079016 | -3.639233 |
| 4 | 28079017 | -3.713322 |
| 5 | 28079018 | -3.731853 |
| 6 | 28079024 | -3.747347 |
| 7 | 28079027 | -3.580031 |
| 8 | 28079035 | -3.703172 |
| 9 | 28079036 | -3.645306 |

|    | id       | lon       |
|----|----------|-----------|
| 10 | 28079038 | -3.707128 |
| 11 | 28079039 | -3.711542 |
| 12 | 28079040 | -3.651522 |
| 13 | 28079047 | -3.686825 |
| 14 | 28079048 | -3.690367 |
| 15 | 28079049 | -3.682583 |
| 16 | 28079050 | -3.688769 |
| 17 | 28079054 | -3.612117 |
| 18 | 28079055 | -3.580747 |
| 19 | 28079056 | -3.718728 |
| 20 | 28079057 | -3.660503 |
| 21 | 28079058 | -3.774611 |
| 22 | 28079059 | -3.609072 |
| 23 | 28079060 | -3.689761 |

## Line chart

```
In [11]: data.plot.line(subplots=True)
```


Out[11]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)



## Line chart

```
In [12]: data.plot.line()
```

Out[12]: <AxesSubplot:>




## Bar chart

```
In [13]: b=data[0:50]
```

```
In [14]: b.plot.bar()
```


```
Out[14]: <AxesSubplot:>
```



## Histogram

```
In [15]: data.plot.hist()
```


```
Out[15]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

```
In [16]: data.plot.area()
```


```
Out[16]: <AxesSubplot:>
```



## Box chart

```
In [17]: data.plot.box()
```


```
Out[17]: <AxesSubplot:>
```



## Pie chart

```
In [19]: b.plot.pie(y='id')
```


Out[19]: <AxesSubplot:ylabel='id'>



## Scatter chart

```
In [106... data.plot.scatter(x='id' ,y='lon')
```

Out[106... <AxesSubplot:xlabel='id', ylabel='lon'>



In [22]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24 entries, 0 to 23
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   id           24 non-null    int64
1   name         24 non-null    object
2   address      24 non-null    object
3   lon          24 non-null    float64
4   lat          24 non-null    float64
5   elevation    24 non-null    int64
dtypes: float64(2), int64(2), object(2)
memory usage: 1.3+ KB
```

In [19]:

df.columns

```
Out[19]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
              'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

In [17]:

df.describe()

Out[17]:

|              | BEN          | CO           | EBE          | NMHC         | NO           | NO_2         |          |
|--------------|--------------|--------------|--------------|--------------|--------------|--------------|----------|
| <b>count</b> | 10916.000000 | 10916.000000 | 10916.000000 | 10916.000000 | 10916.000000 | 10916.000000 | 10916.00 |
| <b>mean</b>  | 0.784014     | 0.279333     | 0.992213     | 0.215755     | 18.795529    | 31.262642    | 44.23    |
| <b>std</b>   | 0.632755     | 0.167922     | 0.804554     | 0.075169     | 40.038872    | 27.234732    | 29.53    |
| <b>min</b>   | 0.100000     | 0.100000     | 0.100000     | 0.050000     | 0.000000     | 1.000000     | 1.00     |
| <b>25%</b>   | 0.400000     | 0.200000     | 0.500000     | 0.160000     | 1.000000     | 9.000000     | 18.00    |
| <b>50%</b>   | 0.600000     | 0.200000     | 0.800000     | 0.220000     | 3.000000     | 24.000000    | 44.00    |
| <b>75%</b>   | 0.900000     | 0.300000     | 1.200000     | 0.250000     | 18.000000    | 47.000000    | 65.00    |
| <b>max</b>   | 7.000000     | 2.500000     | 9.700000     | 0.670000     | 525.000000   | 225.000000   | 157.00   |

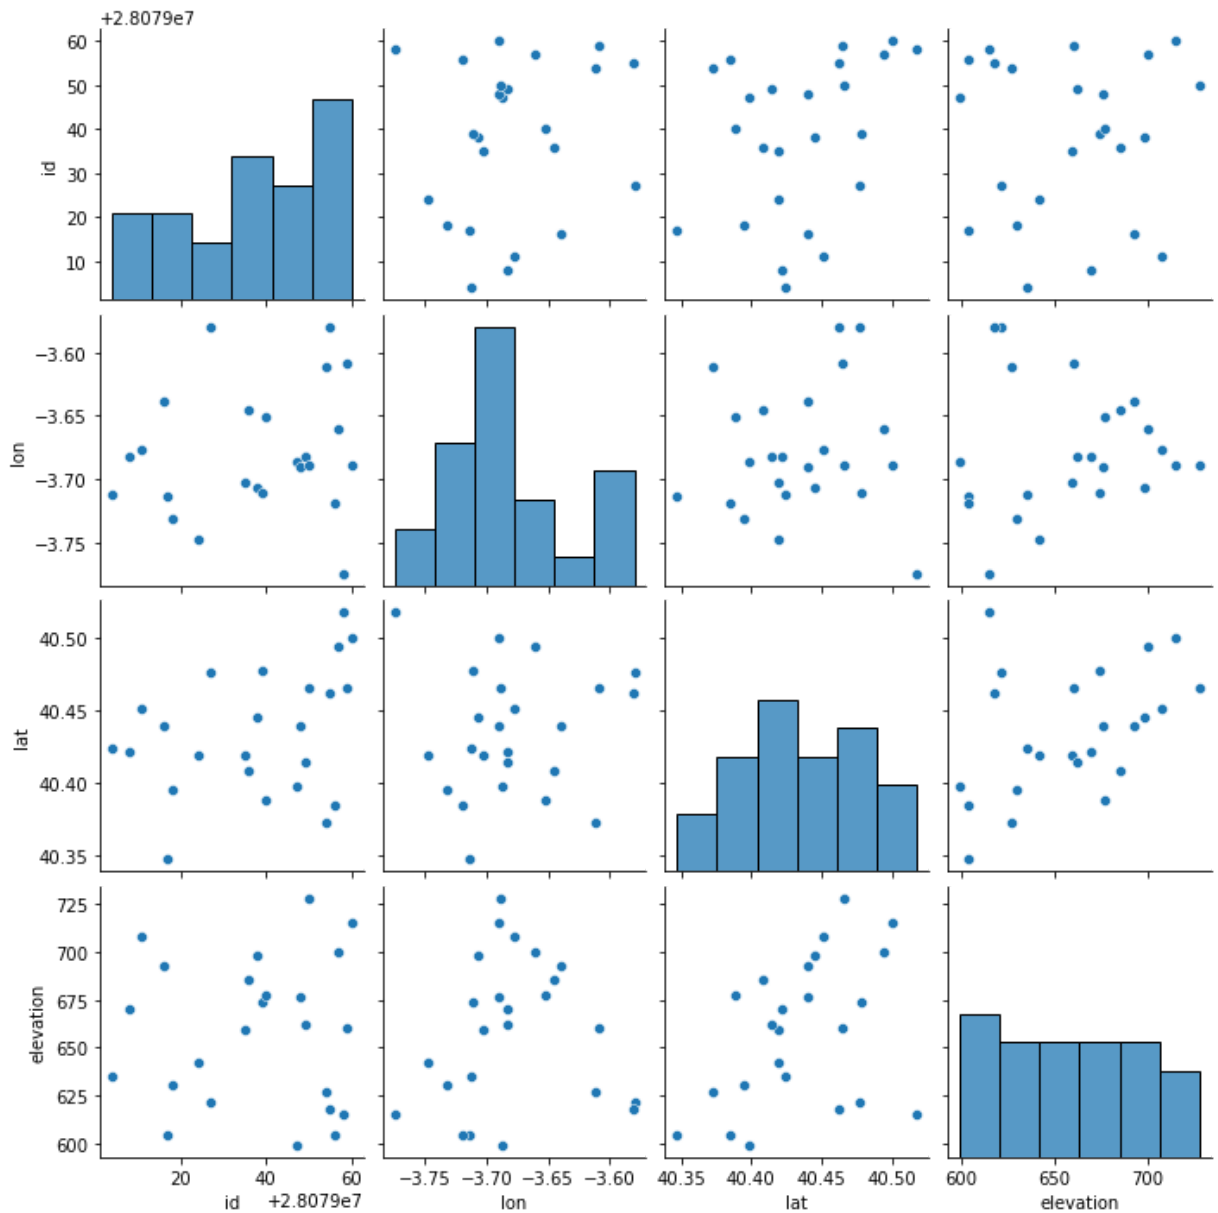
In [23]:

df1=df[['id', 'name', 'address', 'lon', 'lat', 'elevation']]

# EDA AND VISUALIZATION

```
In [24]: sns.pairplot(df1[0:50])
```

```
Out[24]: <seaborn.axisgrid.PairGrid at 0x190514fe400>
```




```
In [26]: sns.distplot(df1['id'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)


```
Out[26]: <AxesSubplot:xlabel='id', ylabel='Density'>
```





In [27]: `sns.heatmap(df1.corr())`

Out[27]: <AxesSubplot:>



In [29]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24 entries, 0 to 23
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id          24 non-null    int64
1   name        24 non-null    object
2   address     24 non-null    object
3   lon         24 non-null    float64
4   lat         24 non-null    float64
5   elevation   24 non-null    int64
dtypes: float64(2), int64(2), object(2)
memory usage: 1.3+ KB
```

## TO TRAIN THE MODEL AND MODEL BUILDING

In [30]: `x=df[['id']]`  
`y=df['elevation']`

```
In [31]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression

```
In [32]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[32]: LinearRegression()

```
In [33]: lr.intercept_
```

Out[33]: 4206047.767358211


```
In [34]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[34]:

|    | Co-efficient |
|----|--------------|
| id | -0.14977     |

```
In [35]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[35]: <matplotlib.collections.PathCollection at 0x190528c32e0>



## ACCURACY

```
In [36]: lr.score(x_test,y_test)
```

Out[36]: -0.25910603716409164

```
In [37]: lr.score(x_train,y_train)
```

Out[37]: 0.004162373438242106

## Ridge and Lasso

```
In [38]: from sklearn.linear_model import Ridge,Lasso
```

```
In [39]: rr=Ridge(alpha=10)
         rr.fit(x_train,y_train)
```

Out[39]: Ridge(alpha=10)

## Accuracy(Ridge)

```
In [40]: rr.score(x_test,y_test)
```

Out[40]: -0.2590389528358372

```
In [41]: rr.score(x_train,y_train)
```

Out[41]: 0.004162352167841776

```
In [42]: la=Lasso(alpha=10)
         la.fit(x_train,y_train)
```

Out[42]: Lasso(alpha=10)

```
In [43]: la.score(x_train,y_train)
```

Out[43]: 0.003918492372358862

## Accuracy(Lasso)

```
In [44]: la.score(x_test,y_test)
```

Out[44]: -0.2524587869827213

## Elastic Net

```
In [45]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(x_train,y_train)
```

Out[45]: ElasticNet()

```
In [46]: en.coef_
```

Out[46]: array([-0.14768952])

In [47]: `en.intercept_`

Out[47]: 4147633.394964873

In [48]: `prediction=en.predict(x_test)`

In [49]: `en.score(x_test,y_test)`

Out[49]: -0.2586953411650055

## Evaluation Metrics

In [50]: 

```
from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

29.966904852713924  
1350.2654272347597  
36.74595797138455

## Logistic Regression

In [58]: `from sklearn.linear_model import LogisticRegression`

In [59]: `feature_matrix=df[['id']]`  
`target_vector=df['elevation']`

In [60]: `feature_matrix.shape`

Out[60]: (24, 1)

In [61]: `target_vector.shape`

Out[61]: (24,)

In [62]: `from sklearn.preprocessing import StandardScaler`

In [63]: `fs=StandardScaler().fit_transform(feature_matrix)`

In [64]: `logr=LogisticRegression(max_iter=10000)`  
`logr.fit(fs,target_vector)`

```
Out[64]: LogisticRegression(max_iter=10000)
```

```
In [70]: observation=[[1]]
```

```
In [71]: prediction=logr.predict(observation)
print(prediction)
```

```
[604]
```

```
In [72]: logr.classes_
```

```
Out[72]: array([599, 604, 615, 618, 621, 627, 630, 635, 642, 659, 660, 662, 670,
        674, 676, 677, 685, 693, 698, 700, 708, 715, 728], dtype=int64)
```

```
In [73]: logr.score(fs,target_vector)
```

```
Out[73]: 0.16666666666666666
```

```
In [74]: logr.predict_proba(observation)[0][0]
```

```
Out[74]: 0.05149080255479361
```

```
In [75]: logr.predict_proba(observation)
```

```
Out[75]: array([[0.0514908 , 0.07628281, 0.06459573, 0.06111161, 0.02846647,
        0.05992926, 0.01982391, 0.00955813, 0.02542849, 0.03721458,
        0.06573175, 0.05391838, 0.01206135, 0.04186834, 0.05270507,
        0.04305319, 0.03836362, 0.0181069 , 0.04069126, 0.06344616,
        0.01416966, 0.06685293, 0.05512959]])
```

## Random Forest

```
In [76]: from sklearn.ensemble import RandomForestClassifier
```

```
In [77]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[77]: RandomForestClassifier()
```

```
In [95]: parameters={'max_depth':[1,2,3,4,5],
        'min_samples_leaf':[5,10,15,20,25],
        'n_estimators':[10,20,30,40,50]
        }
```

```
In [96]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model\_selection\\_split.py:666: UserWarning: The least populated class in y has only 1 members, which is less than n\_s

```
        splits=2.  
        warnings.warn(("The least populated class in y has only %d"  
Out[96]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                    param_grid={'max_depth': [1], 'min_samples_leaf': [5],  
                                'n_estimators': [10]},  
                    scoring='accuracy')
```

```
In [97]: grid_search.best_score_
```

```
Out[97]: 0.0
```

```
In [98]: rfc_best=grid_search.best_estimator_
```

## Conclusion

## Scores

## Linear Regression

```
In [ ]: lr.score(x_test,y_test)
```

```
In [ ]: lr.score(x_train,y_train)
```

## Lasso

```
In [ ]: la.score(x_test,y_test)
```

## Ridge

```
In [83]: rr.score(x_test,y_test)
```

```
Out[83]: -0.2590389528358372
```

```
In [84]: rr.score(x_train,y_train)
```

```
Out[84]: 0.004162352167841776
```

## Elastic Net

```
In [85]: en.score(x_test,y_test)
```

```
Out[85]: -0.2586953411650055
```

# Logistic Regression

```
In [86]: logit.score(fs,target_vector)
```

```
Out[86]: 0.16666666666666666
```

# Random Forest

```
In [87]: grid_search.best_score_
```

```
Out[87]: 0.125
```

From the above data, we can conclude that logistic regression is preferable to other regression types

```
In [ ]:
```